






# GPU COMPUTING WITH OPENACC



# 3 WAYS TO ACCELERATE APPLICATIONS

Applications

Libraries

“Drop-in”  
Acceleration

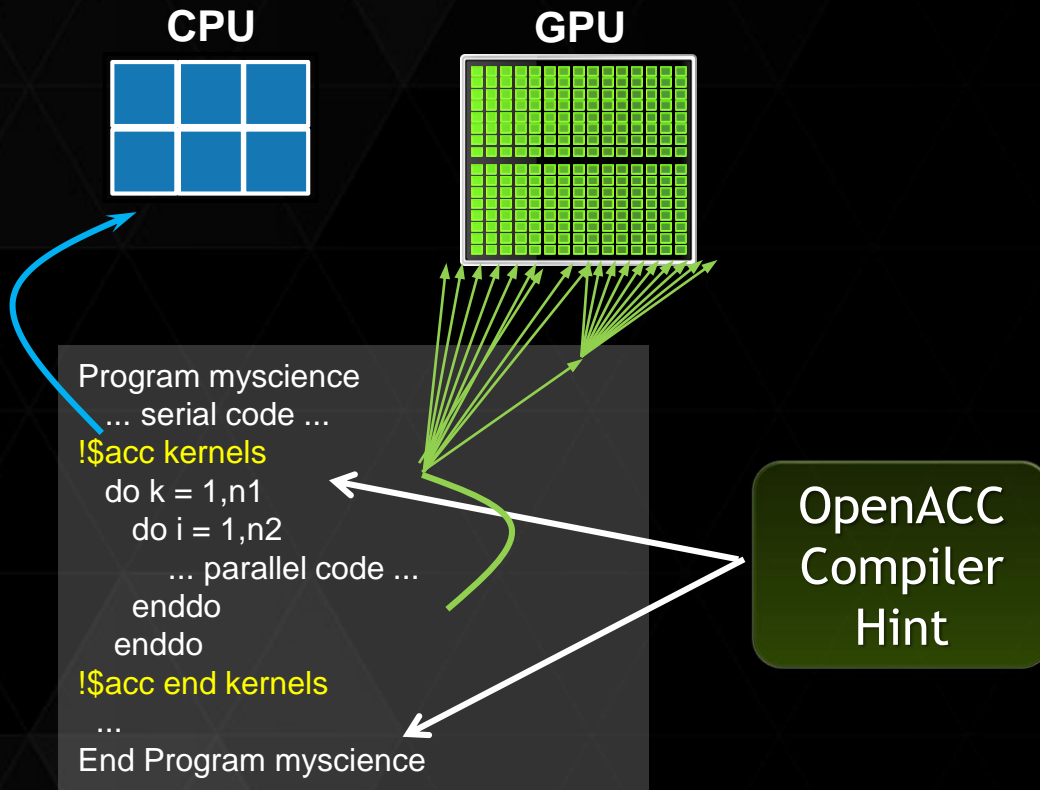
OpenACC  
Directives

Easily Accelerate  
Applications

Programming  
Languages

Maximum  
Flexibility

# OPENACC DIRECTIVES



**Your original  
Fortran or C code**

Simple Compiler hints

Compiler Parallelizes code

Works on many-core GPUs &  
multicore CPUs

# FAMILIAR TO OPENMP PROGRAMMERS

## OpenMP


### CPU




```
main() {  
    double pi = 0.0; long i;  
  
    #pragma omp parallel for reduction(+:pi)  
    for (i=0; i<N; i++)  
    {  
        double t = (double)((i+0.05)/N);  
        pi += 4.0/(1.0+t*t);  
    }  
  
    printf("pi = %f\n", pi/N);  
}
```

## OpenACC

### CPU



### GPU



```
main() {  
    double pi = 0.0; long i;  
  
    #pragma acc kernels  
    for (i=0; i<N; i++)  
    {  
        double t = (double)((i+0.05)/N);  
        pi += 4.0/(1.0+t*t);  
    }  
  
    printf("pi = %f\n", pi/N);  
}
```

# OpenACC Members and Supporters





# DIRECTIVES: EASY & POWERFUL

## Real-Time Object Detection

Global Manufacturer of Navigation Systems



**5x in 40 Hours**

## Valuation of Stock Portfolios using Monte Carlo


Global Technology Consulting Company



**2x in 4 Hours**

## Interaction of Solvents and Biomolecules

University of Texas at San Antonio



**5x in 8 Hours**

“Optimizing code with directives is quite easy, especially compared to CPU threads or writing CUDA kernels. The most important thing is avoiding restructuring of existing code for production applications.”

-- Developer at the Global Manufacturer of Navigation Systems


# FOCUS ON EXPOSING PARALLELISM

With Directives, tuning work focuses on *exposing parallelism*, which makes codes inherently better


## Example: Application tuning work using directives for new Titan system at ORNL

### S3D

Research more efficient combustion with next-generation fuels



- Tuning top 3 kernels (90% of runtime)
- **3 to 6x faster on CPU+GPU vs. CPU+CPU**
- But also improved all-CPU version by 50%



### CAM-SE

Answer questions about specific climate change adaptation and mitigation scenarios

- Tuning top key kernel (50% of runtime)
- **6.5x faster on CPU+GPU vs. CPU+CPU**
- Improved performance of CPU version by 100%

# A VERY SIMPLE EXERCISE: SAXPY

## *SAXPY in C*

```
void saxpy(int n,  
          float a,  
          float *x,  
          float *restrict y)  
{  
    #pragma acc kernels  
    for (int i = 0; i < n; ++i)  
        y[i] = a*x[i] + y[i];  
}  
  
...  
// Perform SAXPY on 1M elements  
saxpy(1<<20, 2.0, x, y);  
...
```

## *SAXPY in Fortran*

```
subroutine saxpy(n, a, x, y)  
    real :: x(:), y(:), a  
    integer :: n, i  
    $!acc kernels  
    do i=1,n  
        y(i) = a*x(i)+y(i)  
    enddo  
    $!acc end kernels  
end subroutine saxpy  
  
...  
$ Perform SAXPY on 1M elements  
call saxpy(2**20, 2.0, x_d, y_d)  
...
```



# DIRECTIVE SYNTAX

- ▶ Fortran

**!\$acc directive [clause [,] clause] ...]**

Often paired with a matching end directive surrounding a structured code block

**!\$acc end directive**

- ▶ C

**#pragma acc directive [clause [,] clause] ...]**

Often followed by a structured code block

# KERNELS: YOUR FIRST OPENACC DIRECTIVE

Each loop executed as a separate *kernel* on the GPU.

```
!$acc kernels
```

```
  do i=1,n  
    a(i) = 0.0  
    b(i) = 1.0  
    c(i) = 2.0  
  end do
```

} kernel 1

```
  do i=1,n  
    a(i) = b(i) + c(i)  
  end do
```

} kernel 2

```
!$acc end kernels
```

## Kernel:

A parallel function  
that runs on the GPU

# KERNELS CONSTRUCT

## Fortran

```
!$acc kernels [clause ...]  
    structured block  
!$acc end kernels
```

## C

```
#pragma acc kernels [clause ...]  
{ structured block }
```

## Clauses

```
if( condition )
```

```
async( expression )
```

Also, any data clause (more later)

# COMPLETE SAXPY EXAMPLE CODE

- ▶ Trivial first example
  - ▶ Apply a loop directive
  - ▶ Learn compiler commands

```
#include <stdlib.h>

void saxpy(int n,
           float a,
           float *x,
           float *restrict y)
{
    #pragma acc kernels
    for (int i = 0; i < n; ++i)
        y[i] = a * x[i] + y[i];
}
```

\*restrict:  
“I promise y does not alias  
x”

```
int main(int argc, char **argv)
{
    int N = 1<<20; // 1 million floats

    if (argc > 1)
        N = atoi(argv[1]);

    float *x = (float*)malloc(N * sizeof(float));
    float *y = (float*)malloc(N * sizeof(float));

    for (int i = 0; i < N; ++i) {
        x[i] = 2.0f;
        y[i] = 1.0f;
    }

    saxpy(N, 3.0f, x, y);

    return 0;
}
```

# COMPILE AND RUN

- ▶ C: `pgcc -acc -ta=nvidia -Minfo=accel -o saxpy_acc saxpy.c`
- ▶ Fortran: `pgf90 -acc -ta=nvidia -Minfo=accel -o saxpy_acc saxpy.f90`
- ▶ Compiler output:

```
pgcc -acc -Minfo=accel -ta=nvidia -o saxpy_acc saxpy.c
saxpy:
```

```
8, Generating copyin(x[:n-1])
```

```
Generating copy(y[:n-1])
```

```
Generating compute capability 1.0 binary
```

```
Generating compute capability 2.0 binary
```

```
9, Loop is parallelizable
```

```
Accelerator kernel generated
```

```
9, #pragma acc loop worker, vector(256) /* blockIdx.x threadIdx.x */
```


```
CC 1.0 : 4 registers; 52 shared, 4 constant, 0 local memory bytes; 100% occupancy
```

```
CC 2.0 : 8 registers; 4 shared, 64 constant, 0 local memory bytes; 100% occupancy
```



# EXAMPLE: JACOBI ITERATION

- ▶ Iteratively converges to correct value (e.g. Temperature), by computing new values at each point from the average of neighboring points.
  - ▶ Common, useful algorithm
  - ▶ Example: Solve Laplace equation in 2D:  $\nabla^2 f(x, y) = 0$



$$A_{k+1}(i, j) = \frac{A_k(i-1, j) + A_k(i+1, j) + A_k(i, j-1) + A_k(i, j+1)}{4}$$

# JACOBI ITERATION C CODE

```
while ( error > tol && iter < iter_max ) {  
    error=0.0;
```

Iterate until converged

```
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {
```

Iterate across matrix  
elements

```
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);
```

Calculate new value from  
neighbors

```
            error = max(error, abs(Anew[j][i] - A[j][i]));
```

Compute max error for  
convergence

```
        }
```

```
    }
```

```
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {
```

```
            A[j][i] = Anew[j][i];
```

```
        }
```

```
    }
```

```
    iter++;
```

```
}
```

Swap input/output arrays

# JACOBI ITERATION FORTRAN CODE

```
do while ( err > tol .and. iter < iter_max )
```



Iterate until converged

```
err=0._fp_kind
```

```
do j=1,m
```



Iterate across matrix elements

```
do i=1,n
```

```
  Anew(i,j) = .25_fp_kind * (A(i+1, j ) + A(i-1, j ) + &  
                             A(i  , j-1) + A(i  , j+1))
```



Calculate new value from neighbors

```
  err = max(err, Anew(i,j) - A(i,j))
```

```
end do
```



Compute max error for convergence

```
end do
```

```
do j=1,m-2
```

```
do i=1,n-2
```

```
  A(i,j) = Anew(i,j)
```

```
end do
```



Swap input/output arrays

```
end do
```

```
iter = iter +1
```

```
end do
```

# EXERCISES

## General instructions (compiling)

- ▶ Exercises are in “exercises/openacc” directory
  - ▶ Solutions in “exercise\_solutions/openacc” directory
- ▶ To compile, use one of the provided makefiles
  - ▶ C: `> make`
  - ▶ Fortran: `> make -f Makefile_f90`
- ▶ Remember these flags
  - ▶ `-acc -ta=nvidia -Minfo=accel`

# EXERCISES

## General instructions (running)

To run, use `sbatch` with one of the provided job files

```
> sbatch runit.acc  
> qstat -u <username>           # prints qsub status
```

Output is placed in `slurm.*` when finished.



# EXERCISE 1

## Jacobi kernels

- ▶ Task: use `acc kernels` to parallelize the Jacobi loop nests
- ▶ Edit `laplace2D.c` or `laplace2D.f90` (your choice)
  - ▶ In the `001-laplace2D-kernels` directory
  - ▶ Add directives where it helps
  - ▶ Figure out the proper compilation flags to use
  - ▶ Optionally: Run OpenMP version with `laplace_omp`
- ▶ Q: can you get a speedup with just kernels directives?
  - ▶ Versus 1 CPU core? Versus 6 CPU cores?

# EXERCISE 1 SOLUTION: OPENACC C

```
while ( error > tol && iter < iter_max ) {  
    error=0.0;  
  
    #pragma acc kernels  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            error = max(error, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    #pragma acc kernels  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
  
    iter++;  
}
```



Execute GPU kernel for  
loop nest



Execute GPU kernel for  
loop nest

# EXERCISE 1 SOLUTION: OPENACC FORTRAN

```
do while ( error > tol .and. iter < iter_max )  
  err=0._fp_kind
```

```
  !$acc kernels
```

```
    do j=1,m  
      do i=1,n
```

```
        Anew(i,j) = 0.25 * (A(i+1,j) + A(i-1,j) + &  
                          A(i,j-1) + A(i,j+1))
```

```
        err = max(err, abs(Anew(i,j) - A(i,j));
```

```
      enddo
```

```
    enddo
```

```
  !$acc end kernels
```

```
  !$acc kernels
```

```
    do j=1, m-2  
      do i=1,n-2
```

```
        A(i,j) = Anew(i,j)
```

```
      enddo
```

```
    enddo
```

```
  !$acc end kernels
```

```
  iter = iter+1
```

```
enddo
```



Execute GPU kernel for  
loop nest



Execute GPU kernel for  
loop nest

# EXERCISE 1: COMPILER OUTPUT (C)

```
pgcc -tp sandybridge-64 -acc -ta=nvidia -Minfo=accel -o laplace2d_acc laplace2d.c
```

```
main:
```

```
56, Generating present_or_copyout(Anew[1:4094][1:4094])
    Generating present_or_copyin(A[:][:])
    Generating Tesla code
57, Loop is parallelizable
59, Loop is parallelizable
    Accelerator kernel generated
    57, #pragma acc loop gang /* blockIdx.y */
    59, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
    63, Max reduction generated for error
68, Generating present_or_copyin(Anew[1:4094][1:4094])
    Generating present_or_copyout(A[1:4094][1:4094])
    Generating Tesla code
69, Loop is parallelizable
71, Loop is parallelizable
    Accelerator kernel generated
    69, #pragma acc loop gang /* blockIdx.y */
    71, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
```

# EXERCISE 1: PERFORMANCE

CPU: Intel E5-2670  
8 Cores @ 2.60 GHz

GPU: NVIDIA Tesla K520m

Execution (4096x4096)	Time (s)	Speedup
CPU 1 OpenMP thread	81.2	--
CPU 2 OpenMP threads	42.4	1.91x
CPU 4 OpenMP threads	25.7	3.16x
CPU 8 OpenMP threads	21.6	5.35x
OpenACC GPU	148.5	0.15x FAIL

Speedup vs. 1 CPU core

Speedup vs. 8 CPU cores



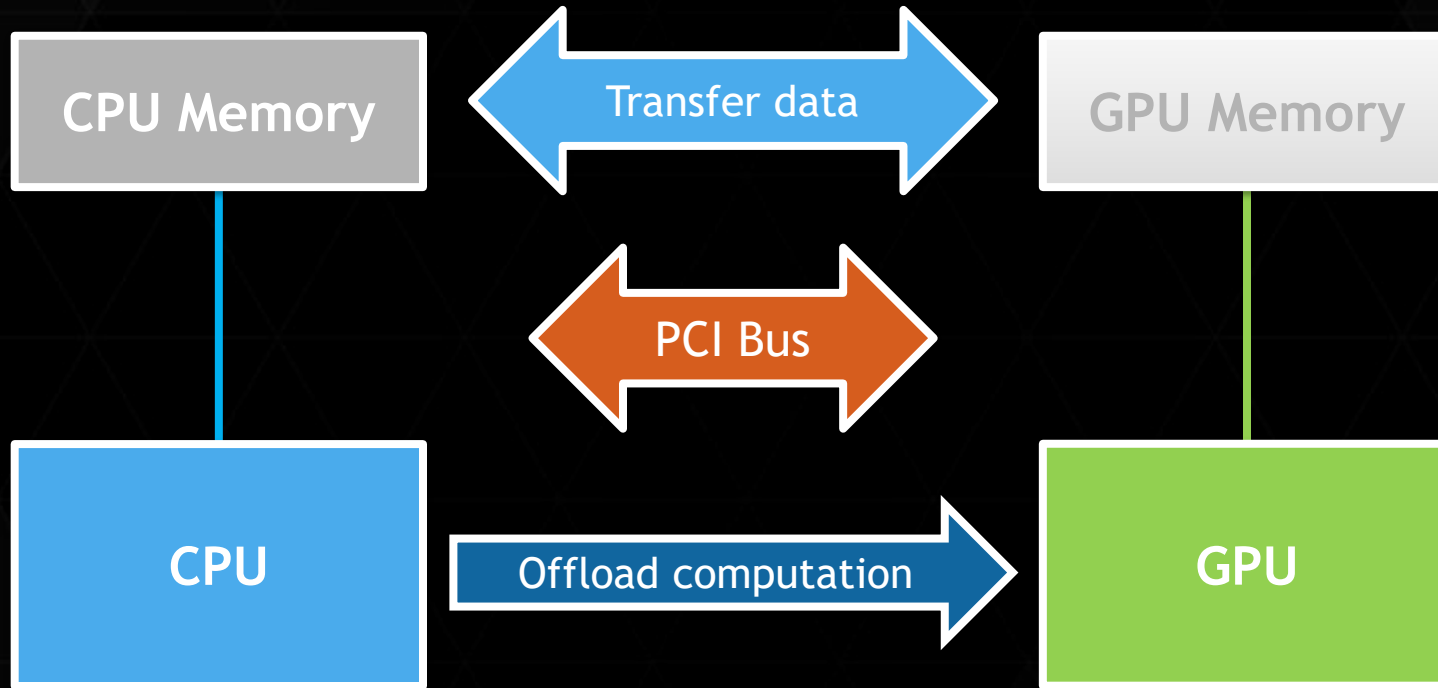
# WHAT WENT WRONG?

PGI\_ACC\_TIME=1

```
time(us): 8,287,436
56: data region reached 1000 times
56: data copyin transfers: 8000
   device time(us): total=182,302 max=52 min=7 avg=22
68: data copyout transfers: 8000
   device time(us): total=226,625 max=118 min=7 avg=28
56: compute region reached 1000 times
59: kernel launched 1000 times
   grid: [32x4094] block: [128]
   device time(us): total=5,000,928 max=5,004 min=4,999 avg=5,000
   elapsed time(us): total=5,045,006 max=6,686 min=5,038 avg=5,045
59: reduction kernel launched 1000 times
   grid: [1] block: [256]
   device time(us): total=239,763 max=241 min=239 avg=239
   elapsed time(us): total=262,790 max=332 min=259 avg=262
68: data region reached 1000 times
68: data copyin transfers: 8000
   device time(us): total=232,372 max=106 min=26 avg=29
77: data copyout transfers: 8000
   device time(us): total=225,256 max=320 min=7 avg=27
68: compute region reached 1000 times
71: kernel launched 1000 times
   grid: [32x4094] block: [128]
   device time(us): total=2,182,750 max=2,186 min=2,178 avg=2,182
   elapsed time(us): total=2,226,348 max=2,458 min=2,220 avg=2,226
```

**Huge Data Transfer Bottleneck!**  
Computation: 10 seconds  
Data malloc/movement: 138

# BASIC CONCEPTS



For efficiency, decouple data movement and compute off-load

# EXCESSIVE DATA TRANSFERS

```
while ( error > tol && iter < iter_max ) {  
    error=0.0;
```

A, Anew resident on host

Copy

**#pragma acc kernels**

A, Anew resident on accelerator

These copies happen  
every iteration of the  
outer while loop!\*

```
for( int j = 1; j < n-1; j++) {  
    for(int i = 1; i < m-1; i++) {  
        Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                             A[j-1][i] + A[j+1][i]);  
        error = max(error, abs(Anew[j][i] - A[j][i]));  
    }  
}
```

A, Anew resident on host

Copy

A, Anew resident on accelerator

}

...

\*Note: there are two #pragma acc kernels, so there are 4 copies per while loop iteration!

The background features a green grid pattern that is warped and curved, creating a sense of depth and movement. Overlaid on this grid are several large, solid black, wavy shapes that resemble liquid droplets or flowing waves, adding a dynamic and organic feel to the design.

# *Data Management*

# DATA CONSTRUCT

## Fortran

```
!$acc data [clause ...]  
    structured block  
!$acc end data
```

## C

```
#pragma acc data [clause ...]  
{ structured block }
```

## General Clauses

```
if( condition )
```

```
async( expression )
```

Manage data movement. Data regions may be nested.



# DATA CLAUSES

`copy ( list )` Allocates memory on GPU and copies data from host to GPU when entering region and copies data to the host when exiting region.

`copyin ( list )` Allocates memory on GPU and copies data from host to GPU when entering region.

`copyout ( list )` Allocates memory on GPU and copies data to the host when exiting region.

`create ( list )` Allocates memory on GPU but does not copy.

`present ( list )` Data is already present on GPU from another containing data region.

`and present_or_copy[in|out], present_or_create, deviceptr.`

# ARRAY SHAPING

- ▶ Compiler sometimes cannot determine size of arrays
  - ▶ Must specify explicitly using data clauses and array “shape”
- ▶ C
  - ▶ `#pragma acc data copyin(a[0:size-1]), copyout(b[s/4:3*s/4])`
- ▶ Fortran
  - ▶ `!$pragma acc data copyin(a(1:size)), copyout(b(s/4:3*s/4))`
- ▶ Note: data clauses can be used on data, kernels or parallel

# EXERCISE 2: JACOBI DATA DIRECTIVES

- ▶ Task: use **acc data** to minimize transfers in the Jacobi example
- ▶ Start from given `laplace2D.c` or `laplace2D.f90` (your choice)
  - ▶ In the **002-laplace2d-data** directory
  - ▶ Add directives where it helps (hint: [do] while loop)
- ▶ Q: What speedup can you get with data + kernels directives?
  - ▶ Versus 6 CPU cores?
  - ▶ `OMP_NUM_THREADS=6 ./laplace2d_omp`

# Exercise 2 Solution: OpenACC C



```
#pragma acc data copy(A), create(Anew)
while ( error > tol && iter < iter_max ) {
    error=0.0;
```

```
#pragma acc kernels
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            error = max(error, abs(Anew[j][i] - A[j][i]));
        }
    }
```

```
#pragma acc kernels
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }

    iter++;
}
```



Copy A in at beginning of loop, out at end. Allocate Anew on accelerator

# Exercise 2 Solution: OpenACC Fortran



```
!$acc data copy(A), create(Anew)
do while ( err > tol .and. iter < iter_max )
    err=0._fp_kind

!$acc kernels
    do j=1,m
        do i=1,n

            Anew(i,j) = .25_fp_kind * (A(i+1, j ) + A(i-1, j ) + &
                                     A(i , j-1) + A(i , j+1))

            err = max(err, Anew(i,j) - A(i,j))
        end do
    end do
!$acc end kernels

...

iter = iter +1
end do
!$acc end data
```



Copy A in at beginning of loop, out at end. Allocate Anew on accelerator

# EXERCISE 2: PERFORMANCE

Execution (4096x4096)	Time (s)	Speedup
CPU 8 OpenMP thread	21.6	--
OpenACC K520m	148.5	0.15x
OpenACC K520m-opt	7.75	2.79x

Speedup vs. 8 CPU cores

# WHAT WENT RIGHT?

```
time(us): 7,423,576
50: data region reached 1 time
50: data copyin transfers: 16
    device time(us): total=305 max=35 min=8 avg=19
82: data copyout transfers: 18
    device time(us): total=182 max=28 min=1 avg=10
56: compute region reached 1000 times
59: kernel launched 1000 times
    grid: [32x4096] block: [128]
    device time(us): total=5,000,403 max=5,003 min=4,998 avg=5,000
    elapsed time(us): total=5,026,969 max=5,194 min=5,020 avg=5,026
59: reduction kernel launched 1000 times
    grid: [1] block: [256]
    device time(us): total=239,865 max=242 min=239 avg=239
    elapsed time(us): total=261,417 max=502 min=258 avg=261
68: compute region reached 1000 times
71: kernel launched 1000 times
    grid: [32x4096] block: [128]
    device time(us): total=2,182,821 max=2,187 min=2,179 avg=2,183
    elapsed time(us): total=2,205,088 max=3,778 min=2,198 avg=2,205
```

**Transfer Bottleneck Eliminated!**  
Computation: 10 seconds  
Data movement: negligible



# FURTHER SPEEDUPS

- ▶ OpenACC gives us more detailed control over parallelization
  - ▶ Via gang, worker, and vector clauses
- ▶ By understanding more about OpenACC execution model and GPU hardware organization, we can get higher speedups on this code
- ▶ By understanding bottlenecks in the code via profiling, we can reorganize the code for higher performance
- ▶ Will tackle these in later exercises

# FINDING PARALLELISM IN YOUR CODE

- ▶ (Nested) for loops are best for parallelization
- ▶ Large loop counts needed to offset GPU/memcpy overhead
- ▶ Iterations of loops must be independent of each other
  - ▶ To help compiler: `restrict` keyword (C), `independent` clause
- ▶ Compiler must be able to figure out sizes of data regions
  - ▶ Can use directives to explicitly control sizes
- ▶ Pointer arithmetic should be avoided if possible
  - ▶ Use subscripted arrays, rather than pointer-indexed arrays.
- ▶ Function calls within accelerated region must be inlineable.

# TIPS AND TRICKS

- (PGI) Use time option to learn where time is being spent
  - -ta=nvidia,time
- Eliminate pointer arithmetic
- Inline function calls in directives regions
  - (PGI): -inline or -inline,levels(<N>)
- Use contiguous memory for multi-dimensional arrays
- Use data regions to avoid excessive memory transfers
- Conditional compilation with `_OPENACC` macro

# PGI 2016 Key New Features

## ■ CPU

- C++14, C++ performance optimizations
- OpenMP and Haswell performance optimizations
- OpenPOWER Alpha release update


## ■ GPU

- App's-driven features and fixes
- OpenACC 2.5 features
- Performance improvements

## ■ Tools

- All-new CPU+GPU performance profiler


# PGI 2016: World-class OpenMP for Haswell



System: 2 x Intel(R) Xeon(R) CPU E5-2698 v3 @ 2.30GHz (32 cores, 64 threads total) 128GB memory  
SPECCompG\_base2012 relative performance as measured by PGI during the weeks of Feb. 1 and Feb. 15, 2016.  
SPECComp® is a registered trademark of the Standard Performance Evaluation Corporation (SPEC).


# PGI 2016: Fortran Performance Improvements

Cycles per element



**Faster Numerical  
Intrinsics  
+  
More SIMD  
Vectorization**

**25% Faster  
WRF Performance**




# PGI 2016: OpenACC Compilers

## Environment




COSMO · NIM · ACME  
SAMI · FV3 · COAMPS

## Chemistry




GAUSSIAN · LS-Dalton  
VWM · NekCEM

## Physics




CloverLeaf · GENE


## CFD



Fun3D · Fluent · HiPSTAR  
INCOMP3D · NUMECA



## SPEC Accel Ratio



Applications-driven  
Features and Tuning

OpenACC 2.5  
Features

11% Faster!

PGI®

System: Intel(R) Core™ i7 3930K CPU @ 3.2 GHz (6 cores total) 16GB memory with NVIDIA Tesla K40c GPU @ 745 MHz. SPECaccel\_acc\_base relative performance as measured by PGI during the weeks of Feb. 15, 2016. SPECaccel™ is a trademark of the Standard Performance Evaluation Corporation (SPEC).



# PGPROF 2016: CPU Profiling, Compiler Feedback

The image shows the PGPROF application window titled "PGPROF (on sb01)". The main pane displays source code from "swim-omp.f". A red circle highlights a section of code between lines 138 and 151. A yellow tooltip box is open over line 140, listing compiler feedback details. The bottom pane shows the "CPU Details" tab with a table of execution events. A red circle highlights the "CPU Details" tab label.

File View Window Help

\*swim-omp.prof swim-omp.f \*swim-acc-data.prof swim-acc-data.f

138 !\$OMP PARALLEL DO  
139 DO 100 1=1, N  
140 Multiple markers at this line  
141 - Generated 6 prefetch instructions for the loop  
142 - Generated vector sse code for the loop  
143 - Generated 5 alternate versions of the loop  
144 - 2 loop-carried redundant expressions removed with 2 operations and 4 arrays  
145 - Intensity = 1.93  
146 V(1;J+1) V(1;J+1)+V(1;J) V(1;J)  
147 100 CONTINUE  
148  
149  
150 C  
151 C PERIODIC CONTINUATION

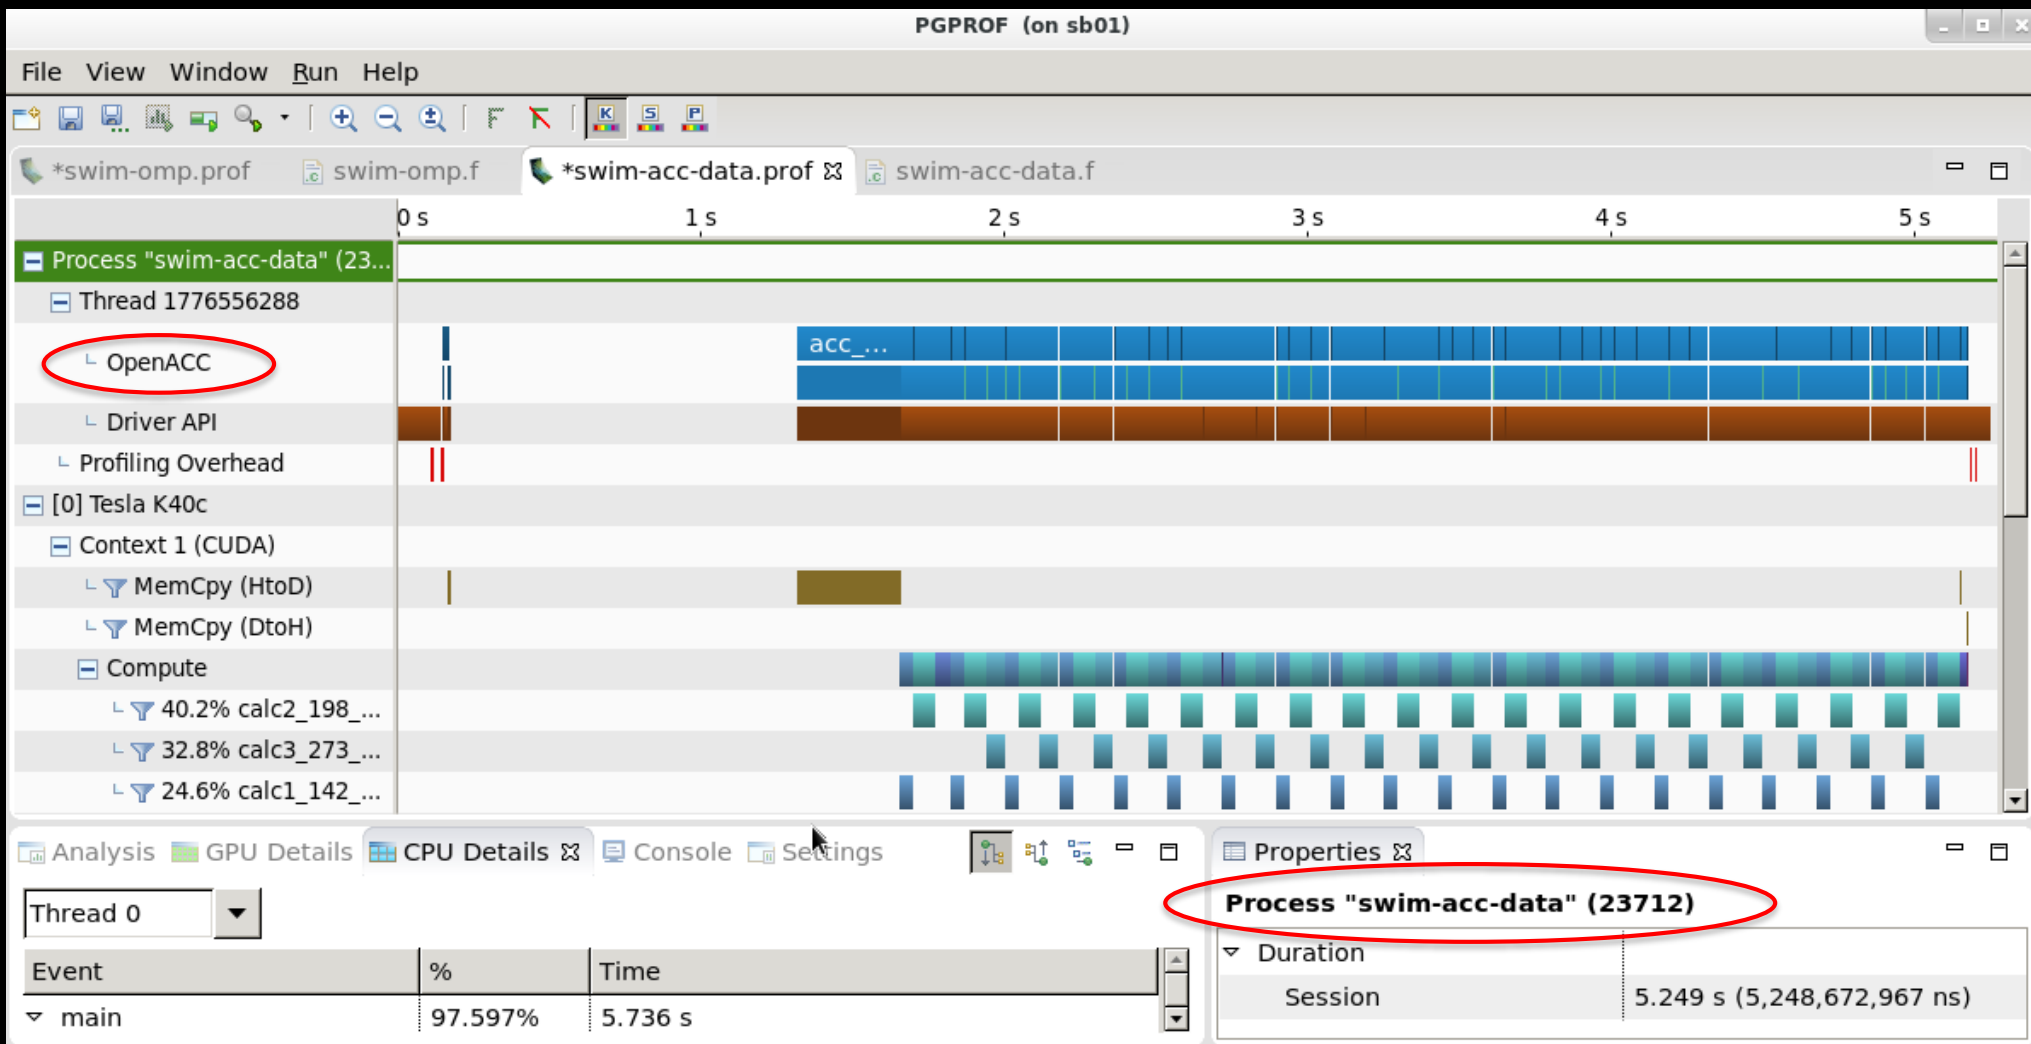
Analysis GPU Details CPU Details Console Settings Properties

Thread 0

Event	%	Time
main	100.124%	8.063 s
MAIN_	100.124%	8.063 s
swim_mod_calc3_	32.174%	2.591 s
swim_mod_calc2_	31.677%	2.551 s
swim mod calc1	26.211%	2.111 s


Select or highlight a single interval to see properties

# PGPROF 2016: OpenACC Profiling



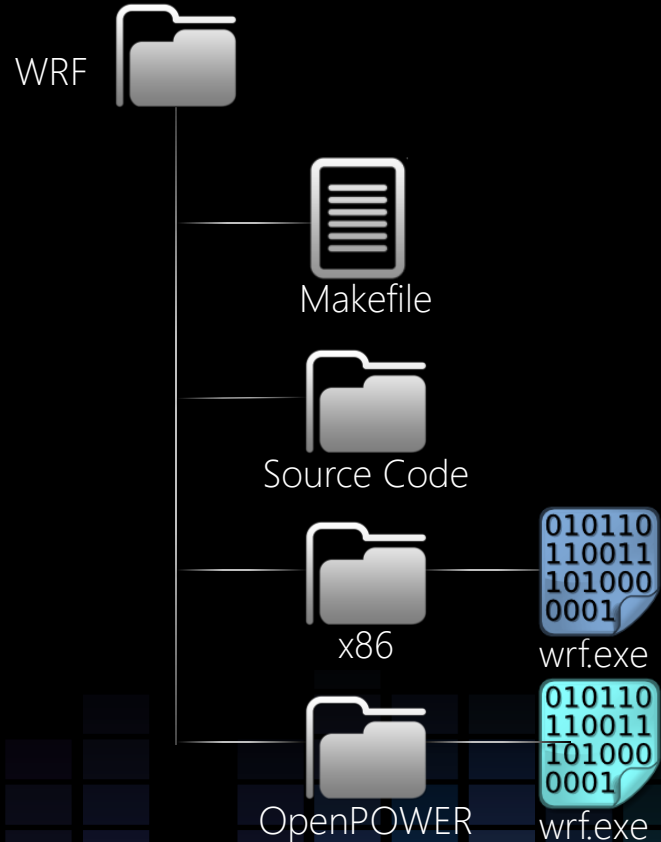
# PGI 16.1-alpha for OpenPOWER+Tesla

- Feature parity with PGI Compilers on Linux/x86+Tesla
- CUDA Fortran, OpenACC, OpenMP, CUDA C/C++ host compiler
- Integrated with IBM's optimized LLVM / OpenPower code generator
- Beta release early in 2016, production later in 2016




# Porting WRF from x86 to OpenPOWER

## Recompile, Run



## WRF 3.5.1 Performance PGI 15.10 vs GNU 5.2



X86 CPU: Intel Xeon E5-2698 v3, 2 sockets, 32 cores total  
POWER CPU: IBM 8247-42L POWER8E, 20 physical cores total  
GNU version 5.2.0; PGI version 15.10

# 2016 OpenACC Toolkit



## PGI Compiler

Free OpenACC compiler for academia



## PGPROF Profiler

Easily find where to add compiler directives



## Code Samples

Learn from examples of real-world algorithms



## Documentation

Quick start guide, Best practices, Forums